

## **Assignment 1 Report**

By Sean Stephen

UOW ID 7311230

Report made on 11/05/2023

### **List of Files :**

Common Files – common.h, common.cpp : Obtained from Tutorials

Task 1 Files – task1\_7311230.cpp (adapted from tutorial1), task1.cl (given as part of assignment 1)

Task 2 Files – task2\_7311230.cpp (adapted from tutorial3), task2.cl (given as part of assignment 1)

Task 3 Files – task3\_7311230.cpp (adapted from tutorial1), arrayVec\_mult.cl (self-written)

### **Explanation + Screenshots of Code**

## common.h

```
1  #pragma once
2  #ifndef _COMMON_H_
3  #define _COMMON_H_
4
5  #define CL_USE_DEPRECATED_OPENCL_2_0_APIS // using OpenCL 1.2, some functions deprecated in OpenCL 2.0
6  #define CL_ENABLE_EXCEPTIONS // enable OpenCL exemptions
7
8  // C++ standard library and STL headers
9  #include <iostream>
10 #include <vector>
11 #include <sstream>
12 #include <fstream>
13
14 // OpenCL header, depending on OS
15 #ifdef __APPLE__
16 #include <OpenCL/cl.hpp>
17 #else
18 #include <CL/cl.hpp>
19 #endif
20
21 // function to handle error
22 void handle_error(cl::Error e);
23
24 // outputs message then quits
25 void quit_program(const std::string str);
26
27 // looks up and displays OpenCL error code as a string
28 const std::string lookup_error_code(cl_int error_code);
29
30 // allows the user to select a device, displays the available platform and device options
31 // returns whether selection was successful, the selected device and its platform
32 bool select_one_device(cl::Platform* platfm, cl::Device* dev);
33
34 // builds program from given filename
35 bool build_program(cl::Program* prog, const cl::Context* ctx, const std::string filename);
36
37 #endif
```

In this file, lines 2 and 3 are to define this header file once only.

Line 5 is to force the system to use OpenCL1.2 instead of the latest version.

Lines 15 – 19 determine if the user's device is MacOS or any other OS such as Windows, as this affects the header file to use as a wrapper for OpenCL C++.

The rest of the lines define the function prototypes for the functions implemented in common.cpp.

### common.cpp

```
// allows the user to select a device, displays the available platform and device options
// returns whether selection was successful, the selected device and its platform
bool select_one_device(cl::Platform* platfm, cl::Device* dev)
{
    std::vector<cl::Platform> platforms; // available platforms
    std::vector< std::vector<cl::Device> > platformDevices; // devices available for each platform
    std::string outputString; // string for output
    unsigned int i, j; // counters

    try {
        // get the number of available OpenCL platforms
        cl::Platform::get(&platforms);
        std::cout << "Number of OpenCL platforms: " << platforms.size() << std::endl;

        // find and store the devices available to each platform
        for (i = 0; i < platforms.size(); i++)
        {
            std::vector<cl::Device> devices; // available devices

            // get all devices available to the platform
            platforms[i].getDevices(CL_DEVICE_TYPE_ALL, &devices);

            // store available devices for the platform
            platformDevices.push_back(devices);
        }

        // display available platforms and devices
        std::cout << "-----" << std::endl;
        std::cout << "Available options:" << std::endl;

        // store options as platform and device indices
        std::vector< std::pair<int, int> > options;
        unsigned int optionCounter = 0; // option counter

        // for all platforms
        for (i = 0; i < platforms.size(); i++)
        {
            // for all devices per platform
            for (j = 0; j < platformDevices[i].size(); j++)
            {
                // display options
                std::cout << "Option " << optionCounter << ": Platform - ";

                // platform vendor name
                outputString = platforms[i].getInfo<CL_PLATFORM_VENDOR>();
                std::cout << outputString << ", Device - ";

                // device name
                outputString = platformDevices[i][j].getInfo<CL_DEVICE_NAME>();
                std::cout << outputString << std::endl;
            }
        }
    }
}
```

The first function implementation is `select_one_device`. It receives a pointer to a platform object, and a device object and returns a Boolean (indicating successful selection of device). Its purpose is to initialise the platform and device objects that the user will choose.

The function uses a try-catch block to catch any OpenCL errors that bubble up.

From the top of the try block, the `platforms` object is first initialised with all the platforms available to the user's system.

Then in the following for-loop, a vector of each platform's devices is added to a vector, where each element corresponds to a platform's list of devices.

The next for-loop is a nested for-loop to display all devices for all platforms.

```

53
54         // store option
55         options.push_back(std::make_pair(i, j));
56         optionCounter++; // increment option counter
57     }
58 }
59
60 std::cout << "\n-----" << std::endl;
61 std::cout << "Select a device: ";
62
63 std::string inputString;
64 unsigned int selectedOption; // option that was selected
65
66 std::getline(std::cin, inputString);
67 std::istringstream stringstream(inputString);
68
69 // check whether valid option selected
70 // check if input was an integer
71 if (stringstream >> selectedOption)
72 {
73     char c;
74
75     // check if there was anything after the integer
76     if (!(stringstream >> c))
77     {
78         // check if valid option range
79         if (selectedOption >= 0 && selectedOption < optionCounter)
80         {
81             // return the platform and device
82             int platformNumber = options[selectedOption].first;
83             int deviceNumber = options[selectedOption].second;
84
85             *platfm = platforms[platformNumber];
86             *dev = platformDevices[platformNumber][deviceNumber];
87
88             return true;
89         }
90     }
91 }
92 // if invalid option selected
93 std::cout << "\n-----" << std::endl;
94 std::cout << "Invalid option." << std::endl;
95 }
96 // catch any OpenCL function errors
97 catch (cl::Error e) {
98     // call function to handle errors
99     handle_error(e);
100 }
101
102 return false;
103 }

```

The user is then prompted to enter which device (that corresponds to its platform) they desire, based on the index(denoted by optionCounter in the nested for-loop, which corresponds to the index of the pairs of indexes i, j, that exist in the vector of vectors of devices which is platformDevices) .

There is some input handling and the index pair is used to initialise the platform and device objects based on the index in the respective vectors.

```

105 // builds program from given filename
106 bool build_program(cl::Program* prog, const cl::Context* ctx, const std::string filename)
107 {
108     // get devices from the context
109     std::vector<cl::Device> contextDevices = ctx->getInfo<CL_CONTEXT_DEVICES>();
110
111     // open input file stream to .cl file
112     std::ifstream programFile(filename);
113
114     // check whether file was opened
115     if (!programFile.is_open())
116     {
117         std::cout << "File not found." << std::endl;
118         return false;
119     }
120
121     // create program string and load contents from the file
122     std::string programString(std::istreambuf_iterator<char>(programFile), (std::istreambuf_iterator<char>()));
123
124     // create program source from one input string
125     cl::Program::Sources source(1, std::make_pair(programString.c_str(), programString.length() + 1));
126     // create program from source
127     *prog = cl::Program(*ctx, source);
128
129     // try to build program
130     try {
131         // build the program for the devices in the context
132         prog->build(contextDevices);
133
134         std::cout << "Program build: Successful" << std::endl;
135         std::cout << "-----" << std::endl;
136     }
137     catch (cl::Error e) {
138         // if failed to build program
139         if (e.err() == CL_BUILD_PROGRAM_FAILURE)
140         {
141             // output program build log
142             std::cout << e.what() << ": Failed to build program." << std::endl;
143
144             // check build status for all all devices in context
145             for (unsigned int i = 0; i < contextDevices.size(); i++)
146             {
147                 // get device's program build status and check for error
148                 // if build error, output build log
149                 if (prog->getBuildInfo<CL_PROGRAM_BUILD_STATUS>(contextDevices[i]) == CL_BUILD_ERROR)
150                 {
151                     // get device name and build log
152                     std::string outputString = contextDevices[i].getInfo<CL_DEVICE_NAME>();
153                     std::string build_log = prog->getBuildInfo<CL_PROGRAM_BUILD_LOG>(contextDevices[i]);
154
155                     std::cout << "Device - " << outputString << ", build log:" << std::endl;
156                     std::cout << "-----" << std::endl;
157                     std::cout << build_log << std::endl;
158                 }
159             }
160             return false;
161         }
162         else
163         {
164             // call function to handle errors
165             handle_error(e);
166         }
167     }
168
169     return true;
170 }
171

```

build\_program is the next function implementation. It receives a program object pointer, a context object pointer and a string for the filename, returning a Boolean to indicate success. The purpose of this function is to build a program (initialise the program object), based on a given OpenCL file.

The function reads the input file into a string and then uses it to create a program source, then create the program with the source and context. Then it tries to build the program for the devices in the current context and any errors are caught and displayed in a verbose manner.

```

172 // function to handle error
173 void handle_error(cl::Error e)
174 {
175     // output OpenCL function that cause the error and the error code
176     std::cout << "Error in: " << e.what() << std::endl;
177     std::cout << "Error code: " << e.err() << " (" << lookup_error_code(e.err()) << ")" << std::endl;
178 }
179

```

handle\_error receives an OpenCL error object and displays the error description and error code.

```

179
180 // function to quit program
181 void quit_program(const std::string str)
182 {
183     std::cout << str << std::endl;
184     std::cout << "Exiting the program..." << std::endl;
185
186     #ifdef _WIN32
187         // wait for a keypress on Windows OS before exiting
188         std::cout << "\npress a key to quit...";
189         std::cin.ignore();
190     #endif
191
192     exit(1);
193 }
194

```

quit\_program is useful as it allows the program to quit early if there are errors in selecting the device or building the program.

```

195 // function to lookup and return error code string
196 const std::string lookup_error_code(cl_int error_code)
197 {
198     // look up error codes as defined in cl.hpp
199     switch (error_code) {
200     case CL_SUCCESS:
201         return "CL_SUCCESS";
202     case CL_DEVICE_NOT_FOUND:
203         return "CL_DEVICE_NOT_FOUND";
204     case CL_DEVICE_NOT_AVAILABLE:
205         return "CL_DEVICE_NOT_AVAILABLE";
206     case CL_COMPILER_NOT_AVAILABLE:
207         return "CL_COMPILER_NOT_AVAILABLE";
208     case CL_MEM_OBJECT_ALLOCATION_FAILURE:
209         return "CL_MEM_OBJECT_ALLOCATION_FAILURE";
210     case CL_OUT_OF_RESOURCES:
211         return "CL_OUT_OF_RESOURCES";
212     case CL_OUT_OF_HOST_MEMORY:
213         return "CL_OUT_OF_HOST_MEMORY";
214     case CL_PROFILING_INFO_NOT_AVAILABLE:
215         return "CL_PROFILING_INFO_NOT_AVAILABLE";
216     case CL_MEM_COPY_OVERLAP:
217         return "CL_MEM_COPY_OVERLAP";
218     case CL_IMAGE_FORMAT_MISMATCH:
219         return "CL_IMAGE_FORMAT_MISMATCH";
220     case CL_IMAGE_FORMAT_NOT_SUPPORTED:
221         return "CL_IMAGE_FORMAT_NOT_SUPPORTED";
222     }
223 }

```

lookup\_error\_code is a very long function that receives an OpenCL integer and matches the correct error code, then returns a string version of it. At the end it also checks if the system is a MacOS system as there are some error codes not defined in MacOS's OpenCL C++ wrapper.

## task1\_7311230.cpp

The first 38 lines are defines, includes and the start of the main method where we declare objects that we need to use later on.

```
38
39 try {
40     // get all available OpenGL platforms
41     cl::Platform::get(&platforms);
42
43     std::cout << "\nThere are " << platforms.size() << " platforms on your system." << std::endl;
44     std::cout << "Here are their information : " << std::endl;
45
46     for (i = 0; i < platforms.size(); i++)
47     {
48         // output the index + 1 for better readability
49         std::cout << "\n\tPlatform " << i + 1 << std::endl;
50
51         // platform name
52         platforms[i].getInfo(CL_PLATFORM_NAME, &outputString);
53         std::cout << "\tPlatform Name : " << outputString << std::endl;
54
55         // get and output platform vendor name
56         outputString = platforms[i].getInfo<CL_PLATFORM_VENDOR>();
57         std::cout << "\tVendor: " << outputString << std::endl;
58
59         // get and output OpenGL version supported by the platform
60         outputString = platforms[i].getInfo<CL_PLATFORM_VERSION>();
61         std::cout << "\tVersion: " << outputString << std::endl;
62     }
63
64     std::cout << "\nPlease input the corresponding platform you would like to select : ";
65
66     std::cin >> platformNum;
67     std::cin.ignore(100, '\n');
68
69     std::cout << "\nYou have selected platform " << platformNum << std::endl;
70
71     // we only want GPU or CPU
72     platforms[platformNum - 1].getDevices(CL_DEVICE_TYPE_CPU | CL_DEVICE_TYPE_GPU, &tempDevices);
73
74     std::cout << "\nThere are " << tempDevices.size() << " devices on your platform." << std::endl;
75     std::cout << "Here are their information : " << std::endl;
76
77     for (i = 0; i < tempDevices.size(); i++)
78     {
79         // output the index + 1 for better readability
80         std::cout << "\n\tDevice " << i + 1 << std::endl;
81
82         // Device name
83         tempDevices[i].getInfo(CL_DEVICE_NAME, &outputString);
84         std::cout << "\tDevice Name : " << outputString << std::endl;
85     }
86
87     std::cout << "\nPlease input the corresponding device you would like to select : ";
88
89     std::cin >> deviceNum;
90     std::cin.ignore(100, '\n');
91
92     std::cout << "\nYou have selected device " << deviceNum << std::endl;
93 }
```

On line 39 we begin our try-catch block. Then we initialise all available platform objects in its vector and display some basic information. Since the first requirement of task 1 is to allow the user to select a device, we must then allow the user to first see all platforms and select one, then select a device under that platform. Once the platform is selected based on its index, we initialise all device objects in its vector and again allow the user to choose which device they want.

```

94     std::cout << "Here is its information : " << std::endl;
95
96
97     // platform name
98     platforms[platformNum - 1].getInfo(CL_PLATFORM_NAME, &outputString);
99     std::cout << "\n\tPlatform Name : " << outputString << std::endl;
100
101     // device type
102     cl_device_type devType;
103     tempDevices[deviceNum - 1].getInfo(CL_DEVICE_TYPE, &devType);
104     if (devType == CL_DEVICE_TYPE_CPU)
105     {
106         outputString = "CL_DEVICE_TYPE_CPU";
107     }
108     else
109     {
110         outputString = "CL_DEVICE_TYPE_GPU";
111     }
112     std::cout << "\tDevice Type : " << outputString << std::endl;
113
114     // device name - returns string
115     tempDevices[deviceNum - 1].getInfo(CL_DEVICE_NAME, &outputString);
116     std::cout << "\tDevice Name : " << outputString << std::endl;
117
118     // number of compute units - returns cl_uint
119     cl_uint maxCompUnits;
120     tempDevices[deviceNum - 1].getInfo(CL_DEVICE_MAX_COMPUTE_UNITS, &maxCompUnits);
121     std::cout << "\tNumber of compute units : " << maxCompUnits << std::endl;
122
123     // max work group size - returns size_t
124     size_t maxWorkGSize;
125     tempDevices[deviceNum - 1].getInfo(CL_DEVICE_MAX_WORK_GROUP_SIZE, &maxWorkGSize);
126     std::cout << "\tMaximum Work Group Size : " << maxWorkGSize << std::endl;
127
128     // max work item dimensions - returns cl_uint
129     cl_uint maxWorkDimensions;
130     tempDevices[deviceNum - 1].getInfo(CL_DEVICE_MAX_WORK_ITEM_DIMENSIONS, &maxWorkDimensions);
131     std::cout << "\tMaximum Work Item Dimensions : " << maxWorkDimensions << std::endl;
132
133
134     // max work item sizes - returns a vector of size_t
135     std::vector<size_t> work_item_sizes;
136     tempDevices[deviceNum - 1].getInfo(CL_DEVICE_MAX_WORK_ITEM_SIZES, &work_item_sizes);
137     std::cout << "\tMaximum Work Item Sizes : ";
138     for (i = 0; i < work_item_sizes.size(); i++)
139     {
140         std::cout << "\n\t\t\t\tDimension " << i + 1 << " : " << work_item_sizes[i];
141     }
142
143     // preferred vector width (int) - returns cl_uint
144     cl_uint prefVecWidth;
145     tempDevices[deviceNum - 1].getInfo(CL_DEVICE_PREFERRED_VECTOR_WIDTH_INT, &prefVecWidth);
146     std::cout << "\n\tPreferred vector width for integers : " << prefVecWidth << std::endl;
147
148     // local memory size - returns cl_ulong
149     cl_ulong mem_size;
150     tempDevices[deviceNum - 1].getInfo(CL_DEVICE_LOCAL_MEM_SIZE, &mem_size);
151     std::cout << "\tLocal Memory Size : " << mem_size << std::endl;
152
153     std::cout << "\n-----\n" << std::endl;

```

Once we have our selected device, we display all needed information, abusing the wrapper `getInfo` method for both the selected platform and device objects. Take note that not all calls to `getInfo` returns strings. Therefore I had to refer to <https://man.opencl.org/clGetDeviceInfo.html> to determine the return type of each flag, then provide the wrapper function with a pointer to the appropriate type. I do realise that the `common.cpp` has an implementation for selecting a device, but I wanted to try it out on my own.



```

154
155 // creating context with device
156 context = cl::Context(tempDevices[deviceNum - 1]);
157
158 // creating command queue with context and device
159 queue = cl::CommandQueue(context, tempDevices[deviceNum - 1]);
160
161 std::cout << "Context and Command Queue created! ~~~" << std::endl;
162
163 // check whether device supports extensions

```

The next requirement was to create a context and a command queue. The context is first created with the selected device, then the command queue is created based on the context and device selected.

```

163 // check whether device supports extensions
164 // returns char[] delimited by spaces. use strtok to read through all, declare if found
165
166 bool fp16 = false, fp64 = false, icd = false;
167
168 char allExt[2000];
169 std::vector<char> allExtensions;
170 tempDevices[deviceNum - 1].getInfo(CL_DEVICE_EXTENSIONS, &allExt);
171
172 // for (i = 0; i < 2000; i++)
173 // {
174 //     std::cout << allExt[i];
175 // }
176
177 // converting to string because I want to use getline() to split the char[] to the different substrings within
178 // tried to use strtok but im not sure why I could not get it to work
179
180 // manually null-terminating allExt to prevent buffer overrun error
181 allExt[2000 - 1] = '\0';
182
183 std::string allExtString(allExt);
184
185 std::stringstream sstream(allExtString);
186
187 std::string tempString;
188
189 std::vector<std::string> allExtStringVec;
190
191 while (std::getline(ss, tempString, ' '))
192 {
193     allExtStringVec.push_back(tempString);
194 }
195
196 // looping through vector of strings to check if the needed extensions are present
197
198 for (i = 0; i < allExtStringVec.size(); i++)
199 {
200     if (allExtStringVec[i] == "cl_khr_fp16")
201     {
202         fp16 = true;
203     }
204     else if (allExtStringVec[i] == "cl_khr_fp64")
205     {
206         fp64 = true;
207     }
208     else if (allExtStringVec[i] == "cl_khr_icd")
209     {
210         icd = true;
211     }
212 }
213
214 std::cout << "\nThe following extensions are / are not supported by the device : " << std::endl;
215 std::cout << "\tcl_khr_fp16 : " << (fp16 ? "Supported" : "Not Supported") << std::endl;
216 std::cout << "\tcl_khr_fp64 : " << (fp64 ? "Supported" : "Not Supported") << std::endl;
217 std::cout << "\tcl_khr_icd : " << (icd ? "Supported" : "Not Supported") << std::endl;
218
219

```

The next requirement was to check if the selected device supported certain extensions. Since the getInfo method required a character array to store the information, but I am not sure of the exact size, I believe 2000 characters is enough to store all the information. Note that I manually null-terminated the char array as my IDE gave a warning that the last character might not be a null-termination, which might cause a buffer overrun.

Since the return of the `getInfo` is a char array where each extension name is delimited by a space, I tried to find ways to split the names into strings. I found that I could use a stringstream object and use `getline` to find delimiters, then store them in a string, which I can add to a vector. I had to initialise the stringstream object with a string therefore I initialised the string with the character array which performs the c-style string to string conversion for me. After all that, I looped through the vector of extension names and checked if the requested extensions were present, assigning true to the respective Booleans if found. Then I displayed if the extensions were supported by the device by using ternary operators with the Booleans.

```
220 // building program
221 if (!build_program(&program, &context, "task1.cl"))
222 {
223     // if OpenCL program build error
224     quit_program("OpenCL program build error.");
225 }
226
227 // display build logs
228 std::string buildLogs = program.getBuildInfo<CL_PROGRAM_BUILD_LOG>(tempDevices[deviceNum - 1]);
229 std::cout << "\nBuild Logs : " << std::endl;
230 std::cout << (buildLogs.size() > 0 ? buildLogs : "No Build Logs were generated.") << std::endl;
231
```

The next requirement was to read the program source code from the provided `task1.cl` file and build the program, displaying the success and build logs (if any). I used the provided function `build_program` in `common.cpp` to build the program with the provided file. Then I used `getBuildInfo` on the created program to initialise a string for the build log and displayed it if there was any.

The last requirement was to find and display the number of kernels in the program, then create the kernels and display all their names. The `createKernels` function creates all kernels in the program at once, as opposed to creating each kernel object manually. This is useful when there are many kernels and the numbers are always changing, removing hard-coding from the equation. From here, the number of kernels is found with the size of the containing vector. We can then loop through each kernel object in the vector and use `getInfo` with the `CL_KERNEL_FUNCTION_NAME` flag to return a string which we can output the kernel's name with.

## Task 1 Result in Terminal

```
There are 1 platforms on your system.
Here are their information :

    Platform 1
    Platform Name : AMD Accelerated Parallel Processing
    Vendor: Advanced Micro Devices, Inc.
    Version: OpenCL 2.1 AMD-APP (3516.0)

Please input the corresponding platform you would like to select : 1

You have selected platform 1

There are 1 devices on your platform.
Here are their information :

    Device 1
    Device Name : gfx1010:xnack-

Please input the corresponding device you would like to select : 1

You have selected device 1
Here is its information :

    Platform Name : AMD Accelerated Parallel Processing
    Device Type : CL_DEVICE_TYPE_GPU
    Device Name : gfx1010:xnack-
    Number of compute units : 20
    Maximum Work Group Size : 256
    Maximum Work Item Dimensions : 3
    Maximum Work Item Sizes :
                                Dimension 1: 1024
                                Dimension 2: 1024
                                Dimension 3: 1024
    Preferred vector width for integers : 1
    Local Memory Size : 65536

-----

Context and Command Queue created! ~~~~

The following extensions are / are not supported by the device :
    cl_khr_fp16 : Supported
    cl_khr_fp64 : Supported
    cl_khr_icd  : Not Supported

Program build: Successful
-----

Build Logs :
C:\Users\User\AppData\Local\Temp\comgr-4e8024\input\CompileSource:29:1: warning: null character ignored [-Wnull-character]
<U+0000>
^
1 warning generated.

Number of Kernels : 5

All Kernel Names :

    add
    copy
    div
    mult
    sub

End of Assignment 1 Task 1

press a key to quit..._
```

## task2\_7311230.cpp

Skipping explanation of all lines until line 48, where it is explained later.

```
45
46 // create and initialise our vector of alphabets
47
48 std::vector<cl_uchar> alphVec;
49
50 // we can make use of the ASCII table to initialise our vector easily
51 // z-a in ASCII is 122 to 97
52 // Z-A in ASCII is 90 to 65
53
54 cl_int i = 0;
55
56 for (i = 122; i >= 97; i--)
57 {
58     alphVec.push_back(i);
59 }
60
61 for (i = 90; i >= 65; i--)
62 {
63     alphVec.push_back(i);
64 }
65
66 // create and initialise our vector of unsigned ints
67
68 std::vector<cl_uint> uIntVec;
69
70 for (i = 0; i <= 1023; i++)
71 {
72     uIntVec.push_back(i);
73 }
74
75
```

Here we fulfil the first requirement by creating and initialising a vector of OpenCL unsigned characters with z-a then Z-A. I am making use of the ASCII values of these characters to initialise the vector with OpenCL integers in a for-loop.

Then I create and initialise the vector of 1024 unsigned integers with a for-loop.

```
76 // creating our three OpenCL buffer objects ,and initialising alphBuffer
77
78 // alphabet buffer
79 cl::Buffer alphBuffer;
80 alphBuffer = cl::Buffer(context, CL_MEM_READ_ONLY | CL_MEM_COPY_HOST_PTR, sizeof(cl_uchar) * alphVec.size(), &alphVec[0]);
81
82 // unsigned char buffer
83 cl::Buffer uCharBuffer;
84 uCharBuffer = cl::Buffer(context, CL_MEM_WRITE_ONLY, sizeof(cl_uchar) * 52);
85
86 // unsigned int buffer
87 cl::Buffer uIntBuffer;
88 uIntBuffer = cl::Buffer(context, CL_MEM_READ_WRITE, sizeof(cl_uint) * 1024);
89
90
```

The next requirement is fulfilled here, by creating buffer objects with the requested flags of read only, write only or read and write. alphBuffer is also created using the copy host pointer flag and initialised with the unsigned character vector we created earlier. Buffers require the read/write flags and the size of the required buffer so that the memory space can be reserved.

```

90
91 // enqueue two OpenCL commands
92
93 // copy first buffer into second buffer
94 queue.enqueueCopyBuffer(alphBuffer, uCharBuffer, 0, 0, sizeof(cl_uchar) * 52);
95
96 queue.enqueueWriteBuffer(uIntBuffer, CL_TRUE, 0, sizeof(cl_uint) * 1024, &uIntVec[0]);
97
98

```

The following requirement is fulfilled by using enqueueCopy/WriteBuffer methods. I am using this as it is simpler than the mapping methods. The copy buffer method requires the source & destination buffers, the offset and frames and the size to copy. The write buffer requires the destination buffer, a flag whether it is blocking or not (blocks the program from proceeding until the write is completed), the offset, size to copy and the source vector.

```

2 cl::Platform platform; // device's platform
3 cl::Device device; // device used
4 cl::Context context; // context for the device
5 cl::Program program; // OpenCL program object
6 cl::Kernel kernel; // a single kernel object
7 cl::CommandQueue queue; // commandqueue for a context and device
8
9
0 try {
1 // select an OpenCL device
2 if (!select_one_device(&platform, &device))
3 {
4 // if no device selected
5 quit_program("Device not selected.");
6 }
7
8 // create a context from device
9 context = cl::Context(device);
10
11
12 // create command queue
13 queue = cl::CommandQueue(context, device);
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99 // build the program
100 if (!build_program(&program, &context, "task2.cl"))
101 {
102 // if OpenCL program build error
103 quit_program("OpenCL program build error.");
104 }
105
106 kernel = cl::Kernel(program, "task2");
107

```

These are the steps (that make use of functions in common.cpp) that allow the user to select the device, create a context and command queue, then create a program and kernel for the given file and kernel names. Note that all objects up till the command queue were created at the start as the queue is needed to enqueue the copy/write buffers, and the buffers need too be created based on a context.

```

107
108     kernel.setArg(0, 12.45f);
109     kernel.setArg(1, uCharBuffer);
110     kernel.setArg(2, uIntBuffer);
111
112     std::cout << "\n\tEnqueue-ing kernel ...\n" << std::endl;
113
114     std::cout << "\t." << std::endl;
115     std::cout << "\t." << std::endl;
116     std::cout << "\t." << std::endl;
117     queue.enqueueTask(kernel);
118
119     std::cout << "\n\tReturned from kernel !!!\n" << std::endl;
120

```

setArg is used to set the arguments provided to the kernel and enqueueTask sends the kernel to be executed.

```

120
121     // read contents from buffers
122
123     std::cout << "Reading the contents of the two buffers now : " << std::endl;
124
125     // two placeholder buffers to read from
126     std::vector<cl_uchar> dispCharBuffer(52);
127     std::vector<cl_uint> dispIntBuffer(1024);
128
129     queue.enqueueReadBuffer(uCharBuffer, CL_TRUE, 0, sizeof(cl_uchar) * 52, &dispCharBuffer[0]);
130
131     std::cout << "\nDone reading Second Buffer. These are its contents in a 4 * 13 matrix :\n" << std::endl;
132
133     for (i = 0; i < 4; i++)
134     {
135         for (int j = 0; j < 13; j++)
136         {
137             std::cout << dispCharBuffer[(j * 1) + (i * 13)] << " ";
138         }
139         std::cout << std::endl;
140     }
141
142     queue.enqueueReadBuffer(uIntBuffer, CL_TRUE, 0, sizeof(cl_uint) * 1024, &dispIntBuffer[0]);
143
144     std::cout << "\nDone reading Third Buffer. These are its contents in a 64 * 16 matrix:\n" << std::endl;
145
146     for (i = 0; i < 64; i++)
147     {
148         for (int j = 0; j < 16; j++)
149         {
150             std::cout << std::setw(4) << dispIntBuffer[(j * 1) + (i * 16)] << " ";
151         }
152         std::cout << std::endl;
153     }
154
155
156
157
158
159 }

```

The contents of the two buffers are read using enqueueReadBuffer, into vectors of the corresponding lengths and types. The method requires a source buffer, CL\_TRUE or CL\_FALSE to indicate blocking, offset, required size to read, as well as the destination vector.

Note that I used a nested for-loop so that I could display the information from the vectors in a 2D format. Accessing the element using  $j + i * \text{row offset}$  also generated a warning from my IDE asking to cast j before adding to i, as they had different sizes. I assume this will be different for different systems.

## Task 2 Result in Terminal

```
Number of OpenCL platforms: 1
-----
Available options:
Option 0: Platform - Advanced Micro Devices, Inc., Device - gfx1010:xnack-
-----
Select a device: 0
Program build: Successful
-----

    Enqueue-ing kernel ...

    .
    .
    .

    Returned from kernel !!!

Reading the contents of the two buffers now :

Done reading Second Buffer. These are its contents in a 4 * 13 matrix :

z y x w v u t s r q p o n
m l k j i h g f e d c b a
Z Y X W V U T S R Q P O N
M L K J I H G F E D C B A

Done reading Third Buffer. These are its contents in a 64 * 16 matrix:

    0   1   2   3   4   5   6   7   8   9   10  11  12  13  14  15
16  17  18  19  20  21  22  23  24  25  26  27  28  29  30  31
32  33  34  35  36  37  38  39  40  41  42  43  44  45  46  47
48  49  50  51  52  53  54  55  56  57  58  59  60  61  62  63
64  65  66  67  68  69  70  71  72  73  74  75  76  77  78  79
80  81  82  83  84  85  86  87  88  89  90  91  92  93  94  95
96  97  98  99 100 101 102 103 104 105 106 107 108 109 110 111
112 113 114 115 116 117 118 119 120 121 122 123 124 125 126 127
128 129 130 131 132 133 134 135 136 137 138 139 140 141 142 143
144 145 146 147 148 149 150 151 152 153 154 155 156 157 158 159
160 161 162 163 164 165 166 167 168 169 170 171 172 173 174 175
176 177 178 179 180 181 182 183 184 185 186 187 188 189 190 191
192 193 194 195 196 197 198 199 200 201 202 203 204 205 206 207
208 209 210 211 212 213 214 215 216 217 218 219 220 221 222 223
224 225 226 227 228 229 230 231 232 233 234 235 236 237 238 239
240 241 242 243 244 245 246 247 248 249 250 251 252 253 254 255
256 257 258 259 260 261 262 263 264 265 266 267 268 269 270 271
272 273 274 275 276 277 278 279 280 281 282 283 284 285 286 287
288 289 290 291 292 293 294 295 296 297 298 299 300 301 302 303
304 305 306 307 308 309 310 311 312 313 314 315 316 317 318 319
320 321 322 323 324 325 326 327 328 329 330 331 332 333 334 335
336 337 338 339 340 341 342 343 344 345 346 347 348 349 350 351
352 353 354 355 356 357 358 359 360 361 362 363 364 365 366 367
368 369 370 371 372 373 374 375 376 377 378 379 380 381 382 383
384 385 386 387 388 389 390 391 392 393 394 395 396 397 398 399
400 401 402 403 404 405 406 407 408 409 410 411 412 413 414 415
416 417 418 419 420 421 422 423 424 425 426 427 428 429 430 431
432 433 434 435 436 437 438 439 440 441 442 443 444 445 446 447
448 449 450 451 452 453 454 455 456 457 458 459 460 461 462 463
464 465 466 467 468 469 470 471 472 473 474 475 476 477 478 479
480 481 482 483 484 485 486 487 488 489 490 491 492 493 494 495
496 497 498 499 500 501 502 503 504 505 506 507 508 509 510 511
512 513 514 515 516 517 518 519 520 521 522 523 524 525 526 527
528 529 530 531 532 533 534 535 536 537 538 539 540 541 542 543
544 545 546 547 548 549 550 551 552 553 554 555 556 557 558 559
560 561 562 563 564 565 566 567 568 569 570 571 572 573 574 575
576 577 578 579 580 581 582 583 584 585 586 587 588 589 590 591
592 593 594 595 596 597 598 599 600 601 602 603 604 605 606 607
608 609 610 611 612 613 614 615 616 617 618 619 620 621 622 623
624 625 626 627 628 629 630 631 632 633 634 635 636 637 638 639
640 641 642 643 644 645 646 647 648 649 650 651 652 653 654 655
656 657 658 659 660 661 662 663 664 665 666 667 668 669 670 671
672 673 674 675 676 677 678 679 680 681 682 683 684 685 686 687
688 689 690 691 692 693 694 695 696 697 698 699 700 701 702 703
704 705 706 707 708 709 710 711 712 713 714 715 716 717 718 719
720 721 722 723 724 725 726 727 728 729 730 731 732 733 734 735
736 737 738 739 740 741 742 743 744 745 746 747 748 749 750 751
752 753 754 755 756 757 758 759 760 761 762 763 764 765 766 767
768 769 770 771 772 773 774 775 776 777 778 779 780 781 782 783
784 785 786 787 788 789 790 791 792 793 794 795 796 797 798 799
800 801 802 803 804 805 806 807 808 809 810 811 812 813 814 815
816 817 818 819 820 821 822 823 824 825 826 827 828 829 830 831
832 833 834 835 836 837 838 839 840 841 842 843 844 845 846 847
848 849 850 851 852 853 854 855 856 857 858 859 860 861 862 863
864 865 866 867 868 869 870 871 872 873 874 875 876 877 878 879
880 881 882 883 884 885 886 887 888 889 890 891 892 893 894 895
896 897 898 899 900 901 902 903 904 905 906 907 908 909 910 911
912 913 914 915 916 917 918 919 920 921 922 923 924 925 926 927
928 929 930 931 932 933 934 935 936 937 938 939 940 941 942 943
944 945 946 947 948 949 950 951 952 953 954 955 956 957 958 959
960 961 962 963 964 965 966 967 968 969 970 971 972 973 974 975
976 977 978 979 980 981 982 983 984 985 986 987 988 989 990 991
992 993 994 995 996 997 998 999 1000 1001 1002 1003 1004 1005 1006 1007
1008 1009 1010 1011 1012 1013 1014 1015 1016 1017 1018 1019 1020 1021 1022 1023

press a key to quit...
```

### task3\_7311230.cpp

```
49 // build the program
50 if(!build_program(&program, &context, "arrayVec_mult.cl"))
51 {
52     // if OpenCL program build error
53     quit_program("OpenCL program build error.");
54 }
55
56 // create a kernel
57 kernel = cl::Kernel(program, "arrayVec_mult");
58
59 // create command queue
60 queue = cl::CommandQueue(context, device);
61
```

Here we build the program, providing the filename of our self-written kernel file, then build a single kernel by specifying the kernel name. All previous lines are standard steps to set-up the

```
61
62 // create buffer
63 addBuffer = cl::Buffer(context, CL_MEM_READ_WRITE | CL_MEM_COPY_HOST_PTR, sizeof(cl_int) * LENGTH, &all3s[0]);
64
65 // get input from user
66
67 int numInput;
68 bool withinRange = false;
69
70 // check if the value is between 2 and 99 inclusive
71 while (!withinRange)
72 {
73     std::cout << "\nDear User, This program takes in a single number between 2 and 99 inclusive. Please enter a number : ";
74     std::cin >> numInput;
75     std::cin.ignore(100, '\n');
76
77     if (numInput >= 2 && numInput <= 99)
78     {
79         withinRange = true;
80     }
81     else
82     {
83         std::cout << "\nNumber not between 2 and 99 Inclusive. Please try again." << std::endl;
84     }
85 }
86
87
88 std::cout << "\nPlease note that you have entered : " << numInput << std::endl;
89
90 std::cout << "\nOriginal Array : \n" << std::endl;
91
92 for (int i = 0; i < 64; i++)
93 {
94     for (int j = 0; j < 8; j++)
95     {
96         std::cout << std::setw(4) << all3s[j + i * 4] << " ";
97     }
98     std::cout << std::endl;
99 }
100
---
```

Since Task 3's requirements is to get a user's input where the number that is input will determine how much to increment the numbers in the 512 element vector, my design is to first create a vector with 512 elements, all initialised with the value 3. The increment starts after 3, therefore conceptually, we need to skip the first element, then add the input number \* index in the vector (after accounting for offset of 1) to 3. This is easily done in a for-loop to get the correct index to multiply, but since we want to use our kernel to speed things up, we need to do it the abovementioned way.

In the above code, I create the buffer where we will read the value of 3, write the result after the multiplication, then read the buffer back into a vector of size 512. Then I read the user's input, making sure it is within the 2-99 inclusive range stated in the requirements, and display the original vector before we enqueue the kernel for execution.



```

101
102 // set kernel arguments
103 kernel.setArg(0, numInput);
104 kernel.setArg(1, addBuffer);
105
106
107 // enqueue kernel for execution
108 //queue.enqueueTask(kernel);
109
110 // offset 1 because we want the first element to remain 1
111 // globalSize 512 because we want to use the index of the array to compute how much to increment
112 cl::NDRange offset(1);
113 cl::NDRange globalSize(512);
114
115 queue.enqueueNDRangeKernel(kernel, offset, globalSize);
116
117 std::cout << "\nKernel enqueued." << std::endl;
118 std::cout << "-----" << std::endl;
119
120 // enqueue command to read from device to host memory
121 queue.enqueueReadBuffer(addBuffer, CL_TRUE, 0, sizeof(cl_int) * LENGTH, &resultVec[0]);
122
123
124 std::cout << "\n Our Results : \n" << std::endl;
125
126 // check the results
127 for (int i = 0; i < 64; i++)
128 {
129     for (int j = 0; j < 8; j++)
130     {
131         std::cout << std::setw(4) << resultVec[j + i * 4] << " ";
132     }
133     std::cout << std::endl;
134 }
135
136 // catch any OpenCL function errors
137 catch (cl::Error e) {
138     // call function to handle errors
139     handle_error(e);
140 }

```

Here, I first set the arguments for the kernel. Then I use enqueueNDRangeKernel to enqueue the kernel for execution. I can make use of the offset of 1 to skip the first element (the first 3). The globalSize is set to 512 so that in the kernel I can use the global id to determine the index of the element that is being accessed in that instance.

enqueueReadBuffer is used to read the contents of the buffer back into a vector, where it is displayed in a 2D format using a nested for-loop.

### arrayVec\_mult.cl

```
1  __kernel void arrayVec_mult(int number,  
2                                __global int* vector) {  
3  
4  
5      int arrayIndex = get_global_id(0);  
6      int arrayValue = vector[arrayIndex];  
7      int incrementValue = number * (arrayIndex);  
8      arrayValue += incrementValue;  
9      vector[arrayIndex] = arrayValue;  
10  
11 }  
12  
13
```

This is the kernel code used in Task 3. The parameter is the integer that the user entered used for our incrementing, as well as the buffer where our elements are stored.

The global id is used to determine which index of the element we are accessing in that instance of the kernel. From there we can first get the original value of 3, determine the amount to increment by using the user entered number to multiply with the index. Then we add the original value of 3 with the increment and reassign that element with the result value.

### Task 3 Result in Terminal

[illegible]

Kernel enqueued.

-----

Our Results :

3	6	9	12	15	18	21	24
15	18	21	24	27	30	33	36
27	30	33	36	39	42	45	48
39	42	45	48	51	54	57	60
51	54	57	60	63	66	69	72
63	66	69	72	75	78	81	84
75	78	81	84	87	90	93	96
87	90	93	96	99	102	105	108
99	102	105	108	111	114	117	120
111	114	117	120	123	126	129	132
123	126	129	132	135	138	141	144
135	138	141	144	147	150	153	156
147	150	153	156	159	162	165	168
159	162	165	168	171	174	177	180
171	174	177	180	183	186	189	192
183	186	189	192	195	198	201	204
195	198	201	204	207	210	213	216
207	210	213	216	219	222	225	228
219	222	225	228	231	234	237	240
231	234	237	240	243	246	249	252
243	246	249	252	255	258	261	264
255	258	261	264	267	270	273	276
267	270	273	276	279	282	285	288
279	282	285	288	291	294	297	300
291	294	297	300	303	306	309	312
303	306	309	312	315	318	321	324
315	318	321	324	327	330	333	336
327	330	333	336	339	342	345	348
339	342	345	348	351	354	357	360
351	354	357	360	363	366	369	372
363	366	369	372	375	378	381	384
375	378	381	384	387	390	393	396
387	390	393	396	399	402	405	408
399	402	405	408	411	414	417	420
411	414	417	420	423	426	429	432
423	426	429	432	435	438	441	444
435	438	441	444	447	450	453	456
447	450	453	456	459	462	465	468
459	462	465	468	471	474	477	480
471	474	477	480	483	486	489	492
483	486	489	492	495	498	501	504
495	498	501	504	507	510	513	516
507	510	513	516	519	522	525	528
519	522	525	528	531	534	537	540
531	534	537	540	543	546	549	552
543	546	549	552	555	558	561	564
555	558	561	564	567	570	573	576
567	570	573	576	579	582	585	588
579	582	585	588	591	594	597	600
591	594	597	600	603	606	609	612
603	606	609	612	615	618	621	624
615	618	621	624	627	630	633	636
627	630	633	636	639	642	645	648
639	642	645	648	651	654	657	660
651	654	657	660	663	666	669	672
663	666	669	672	675	678	681	684
675	678	681	684	687	690	693	696
687	690	693	696	699	702	705	708
699	702	705	708	711	714	717	720
711	714	717	720	723	726	729	732
723	726	729	732	735	738	741	744
735	738	741	744	747	750	753	756
747	750	753	756	759	762	765	768
759	762	765	768	771	774	777	780

press a key to quit...